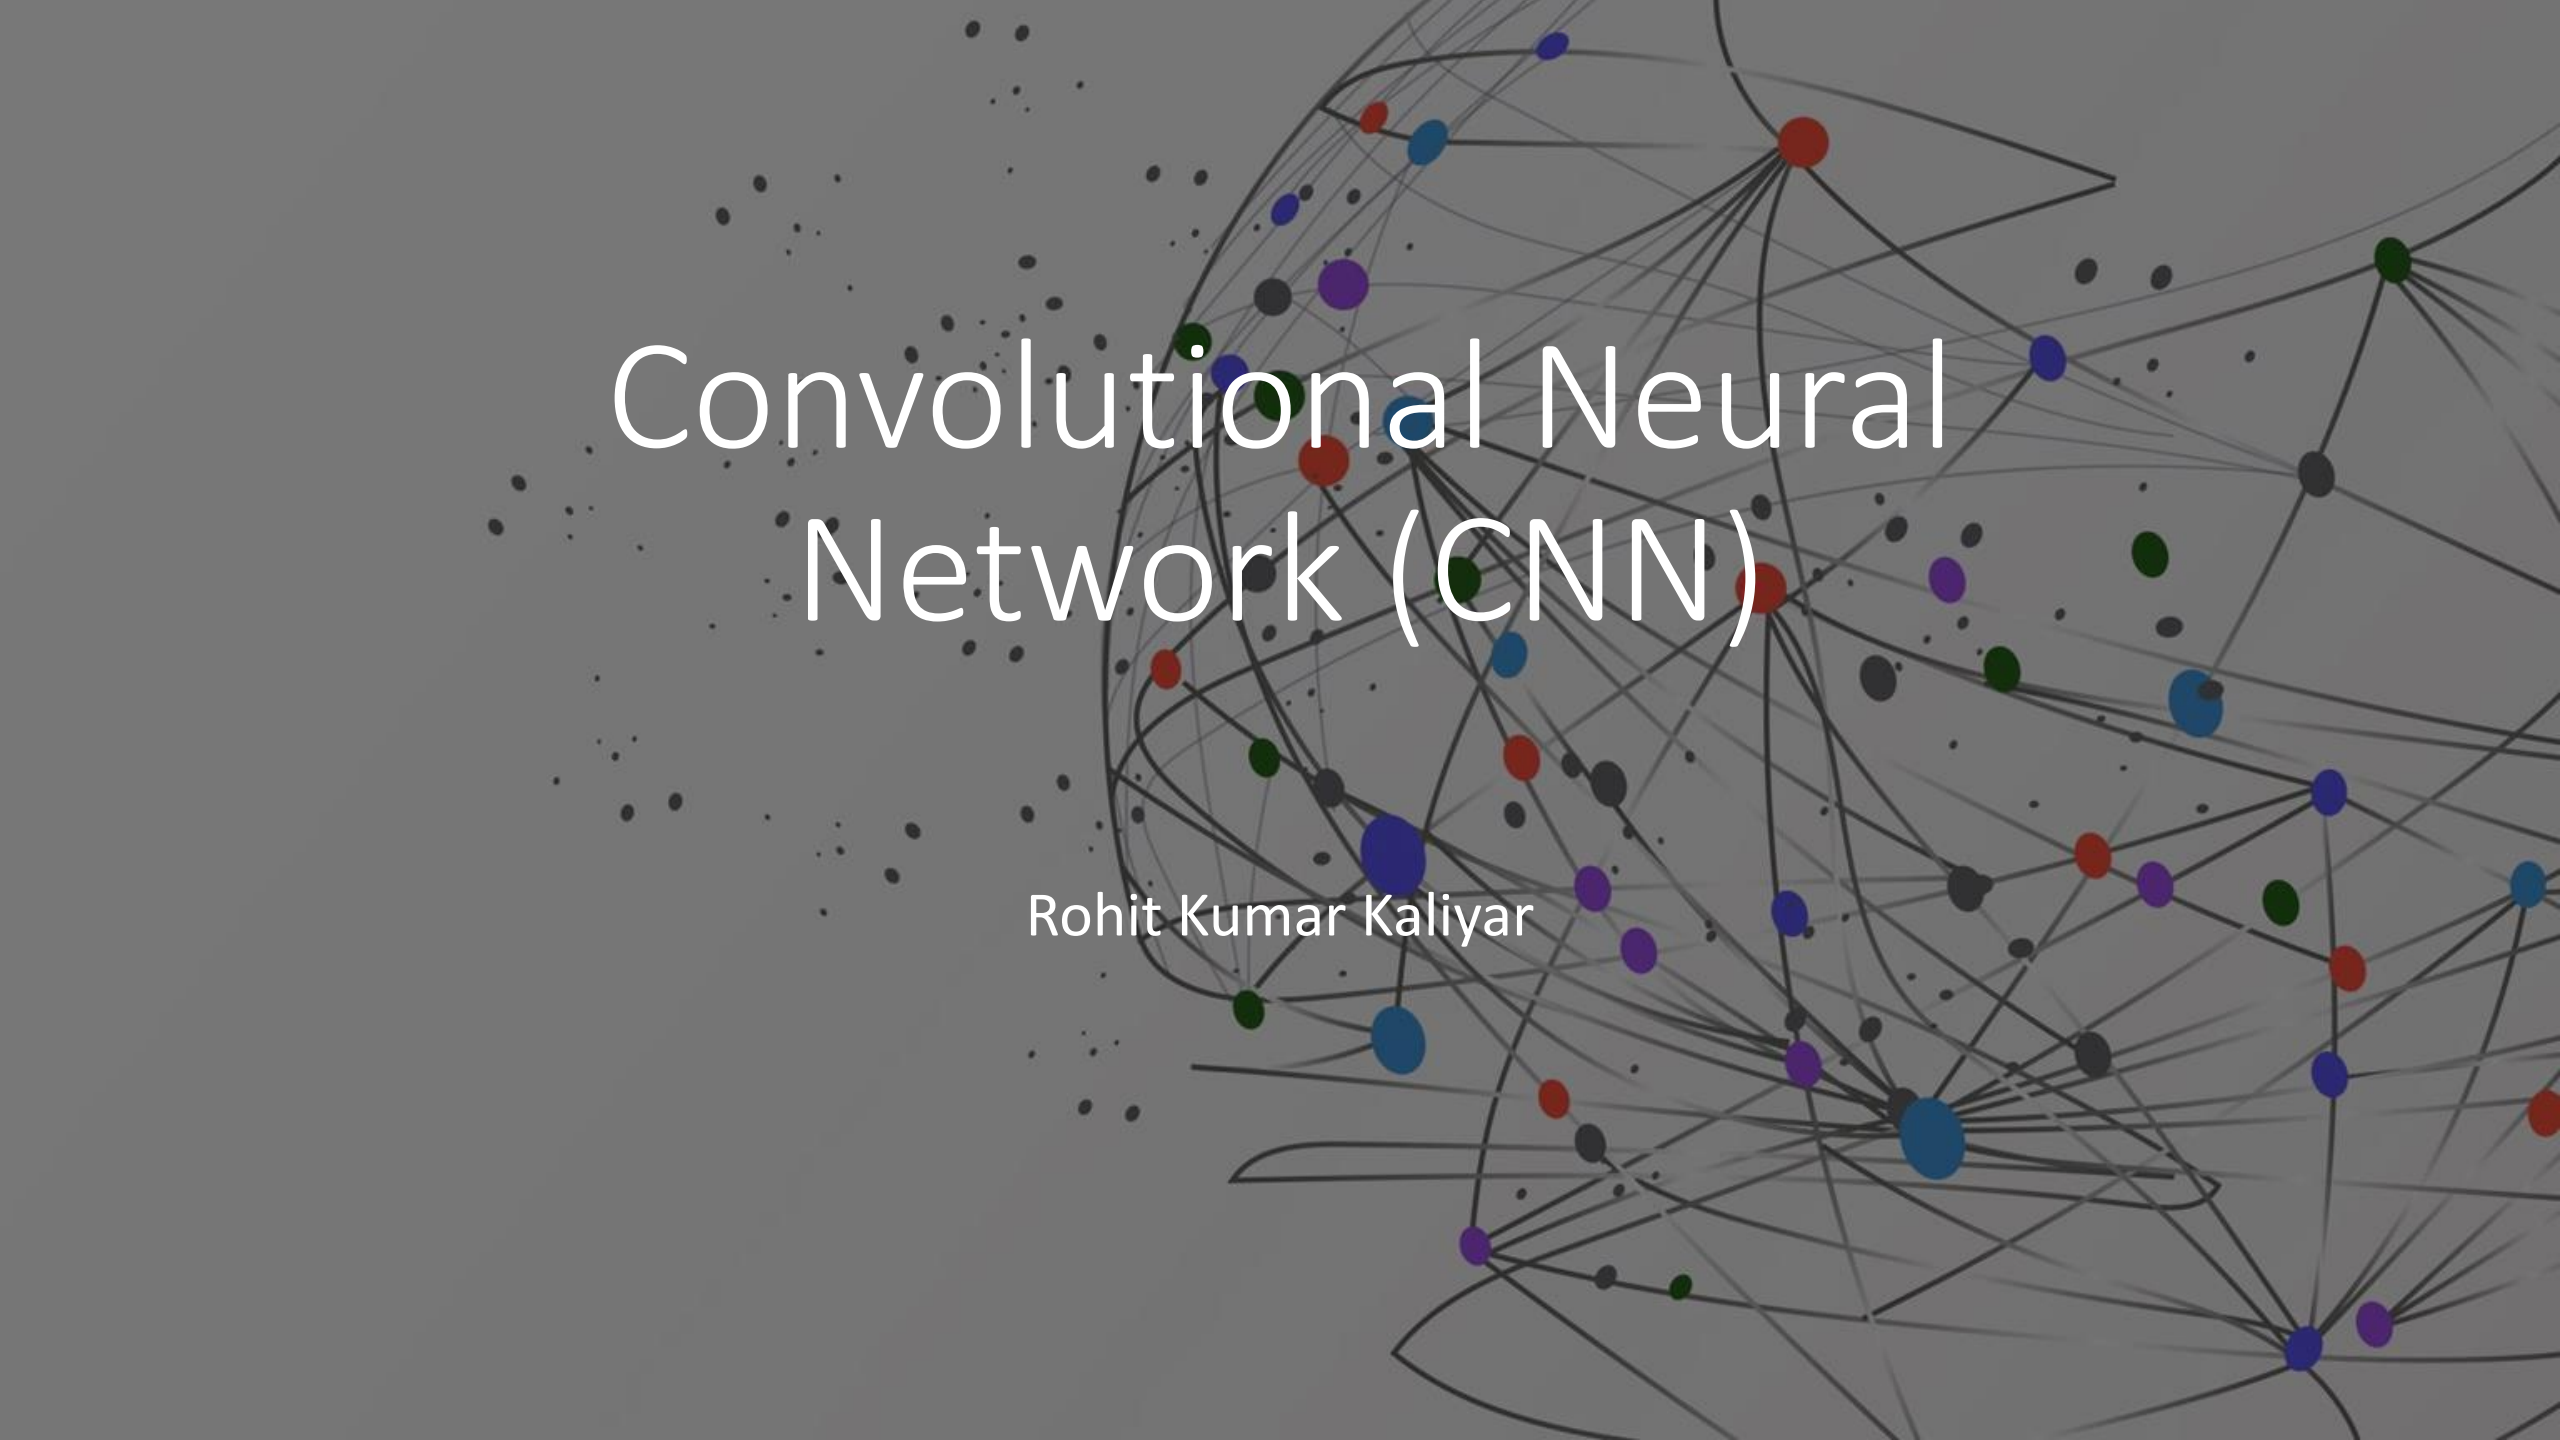



Convolutional Neural Network (CNN)

The background of the slide is a dark gray with a complex, abstract network of thin, light gray lines. These lines connect various colored dots (nodes) in shades of red, blue, green, purple, and black. The nodes are scattered across the frame, with some appearing in small clusters and others in isolation. The overall effect is a sense of interconnectedness and complexity, reminiscent of a neural network or a data visualization.

Rohit Kumar Kaliyar

Why CNN for Image?



Can the network be simplified by considering the properties of images?

Why CNN for Image

- Some patterns are much smaller than the whole image

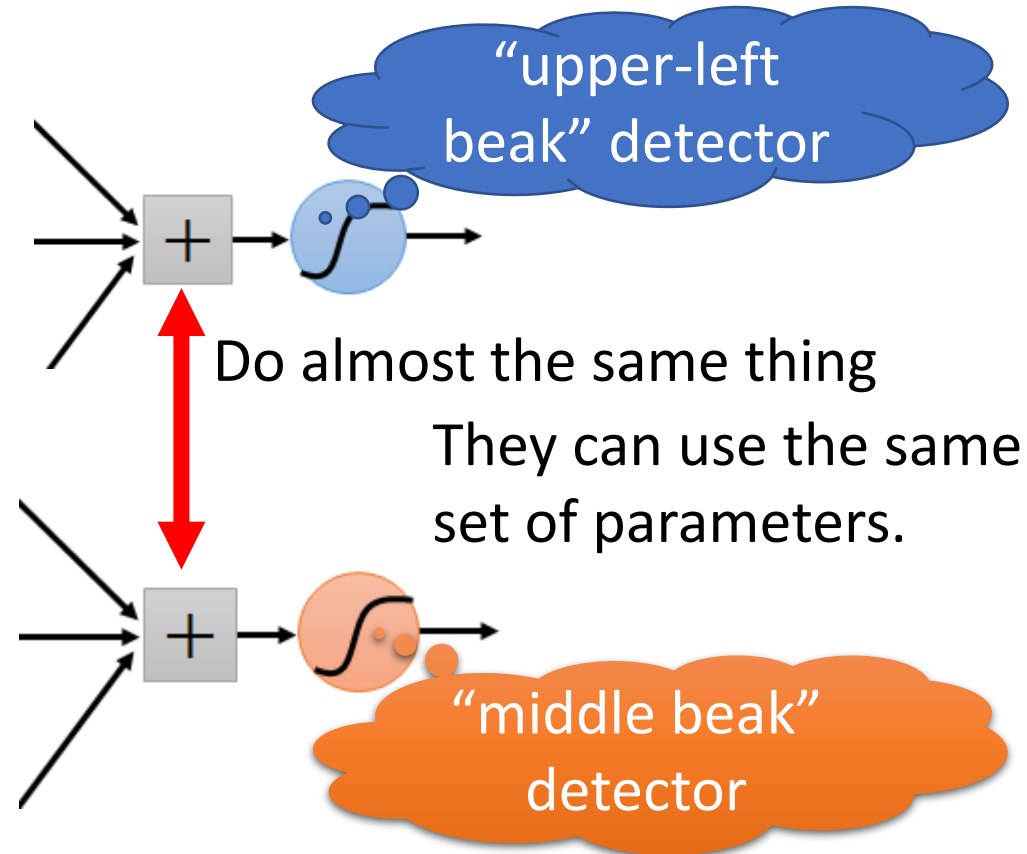

A neuron does not have to see the whole image to discover the pattern.

Connecting to small region with less parameters



Why CNN for Image

- The same patterns appear in different regions.



Why CNN for Image

- Subsampling the pixels will not change the object

bird



subsampling

bird

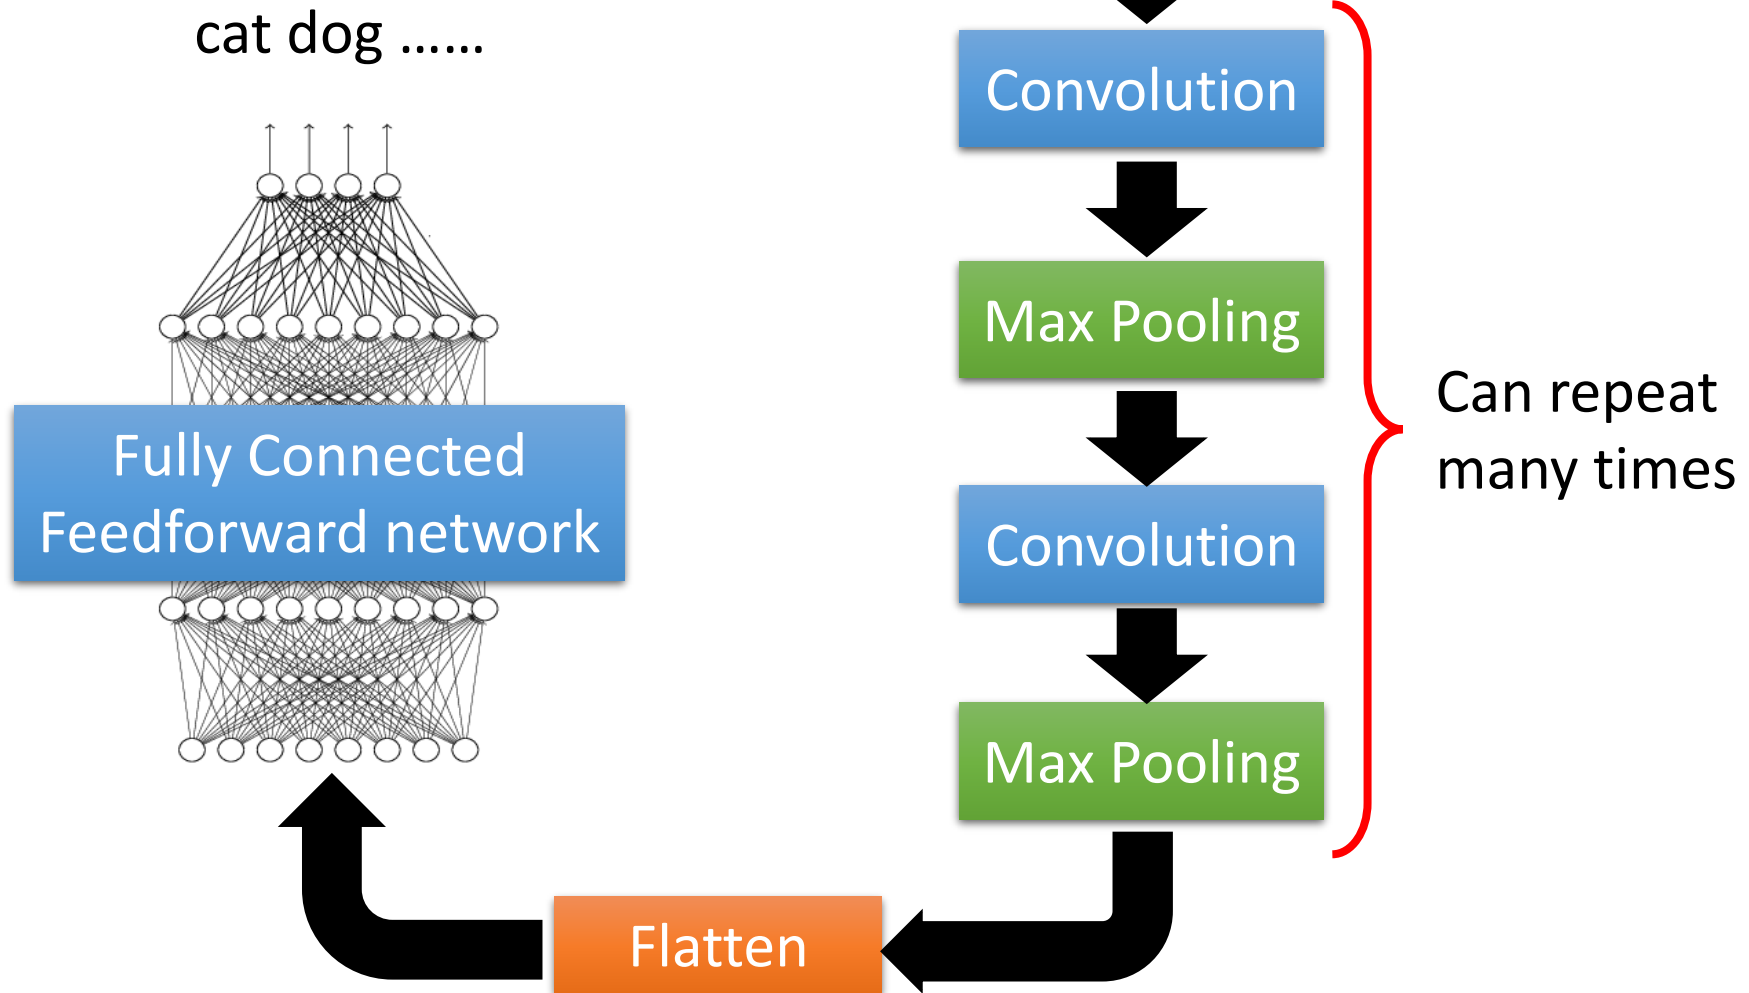


We can subsample the pixels to make image smaller

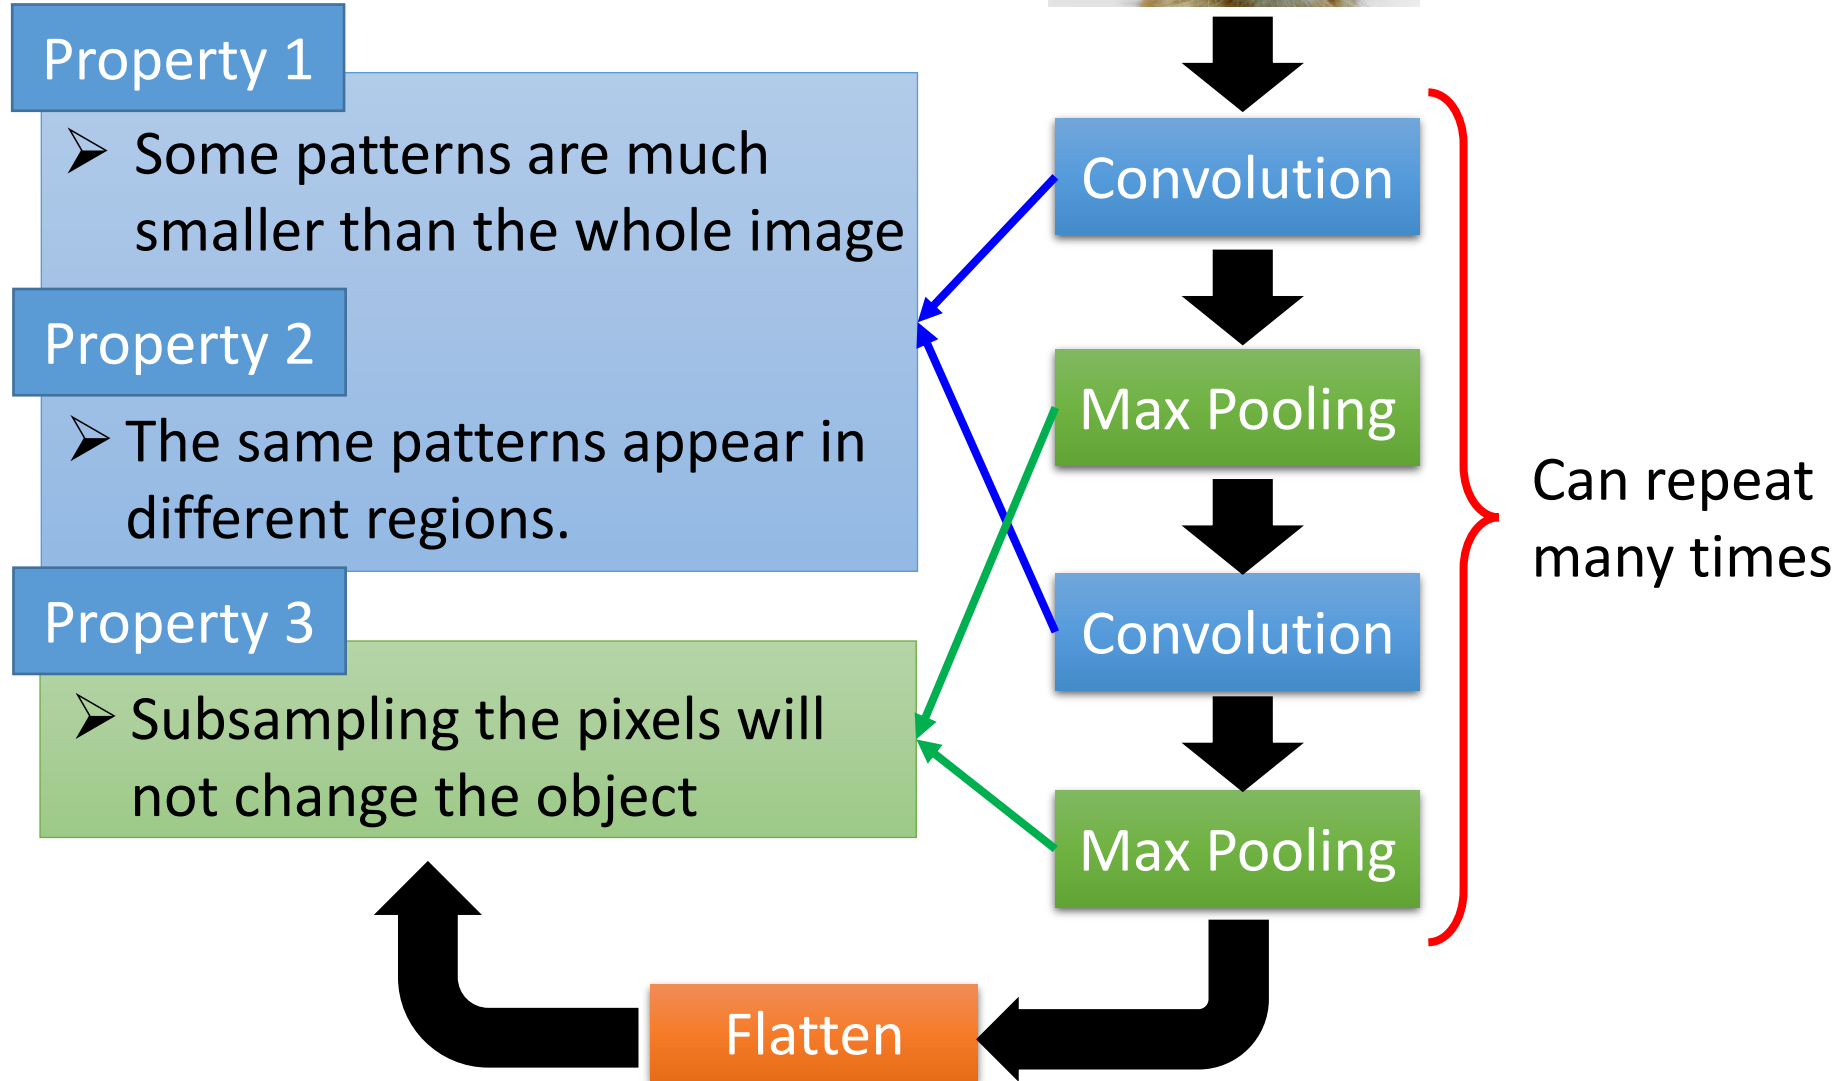


Less parameters for the network to process the image

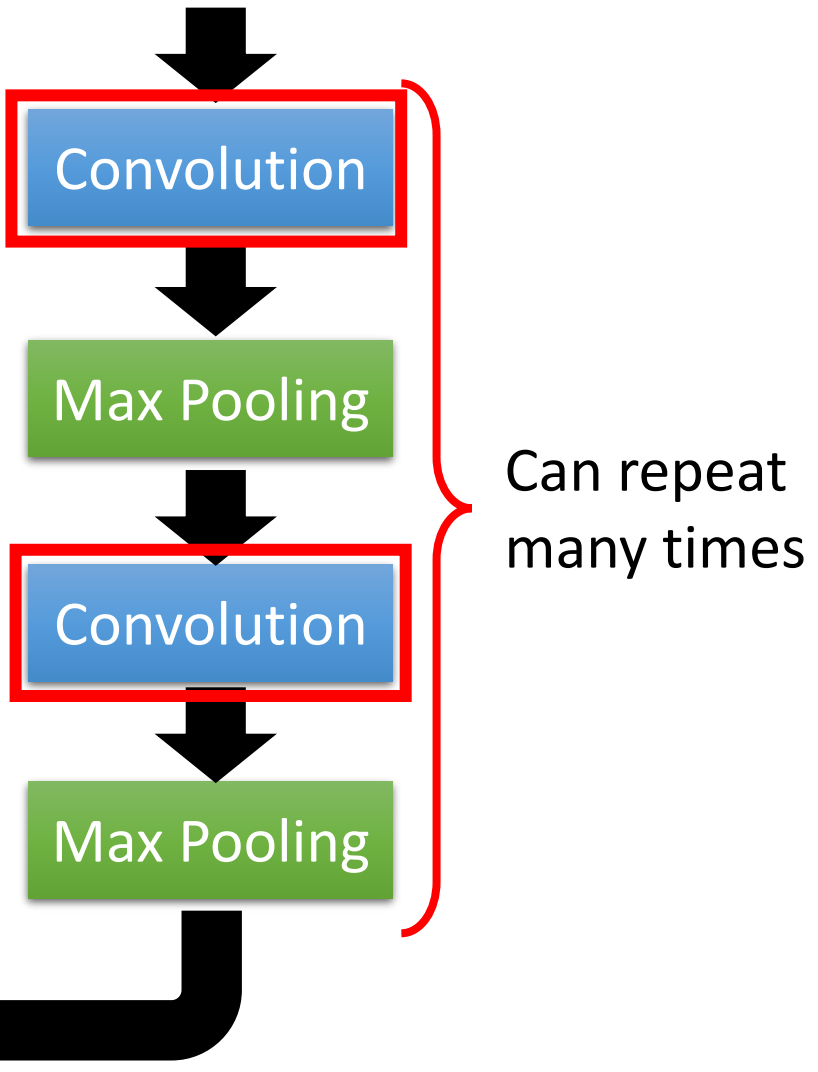

The whole CNN



The whole CNN



The whole CNN



CNN – Convolution

1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
1	0	0	0	1	0
0	1	0	0	1	0
0	0	1	0	1	0

6 x 6 image

Those are the network
parameters to be learned.

1	-1	-1
-1	1	-1
-1	-1	1

Filter 1

Matrix

-1	1	-1
-1	1	-1
-1	1	-1

Filter 2

Matrix

⋮

Property 1

Each filter detects a small
pattern (3 x 3).

CNN – Convolution


stride=1

1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
1	0	0	0	1	0
0	1	0	0	1	0
0	0	1	0	1	0

6 x 6 image

1	-1	-1
-1	1	-1
-1	-1	1

Filter 1



CNN – Convolution


If stride=2

1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
1	0	0	0	1	0
0	1	0	0	1	0
0	0	1	0	1	0

6 x 6 image

1	-1	-1
-1	1	-1
-1	-1	1

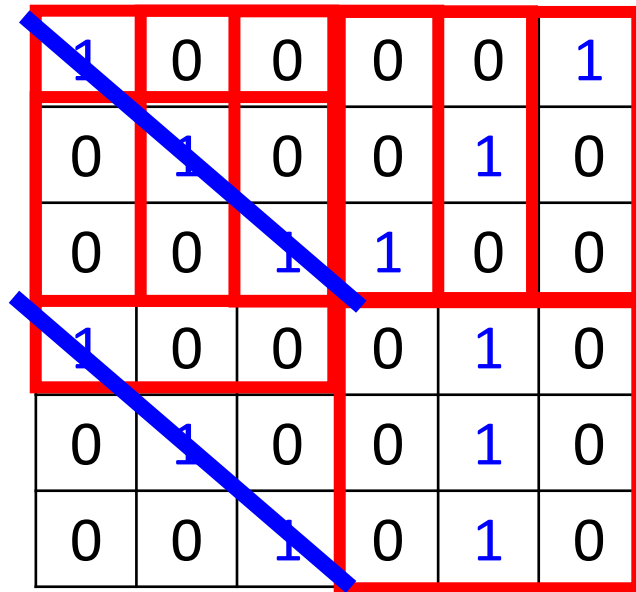
Filter 1



We set stride=1 below

CNN – Convolution

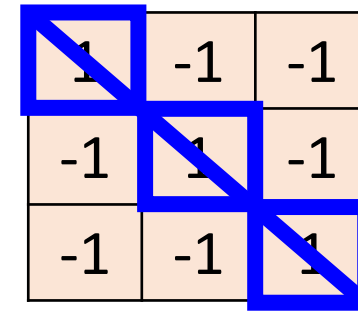
stride=1



A 6x6 grid of numbers representing an image. The values are: Row 1: 1, 0, 0, 0, 0, 1; Row 2: 0, 1, 0, 0, 1, 0; Row 3: 0, 0, 1, 1, 0, 0; Row 4: 1, 0, 0, 0, 1, 0; Row 5: 0, 1, 0, 0, 1, 0; Row 6: 0, 0, 1, 0, 1, 0. A 3x3 red bounding box highlights the top-left corner (rows 1-3, columns 1-3). A blue diagonal line runs from the top-left to the bottom-right of the entire grid.

1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
1	0	0	0	1	0
0	1	0	0	1	0
0	0	1	0	1	0

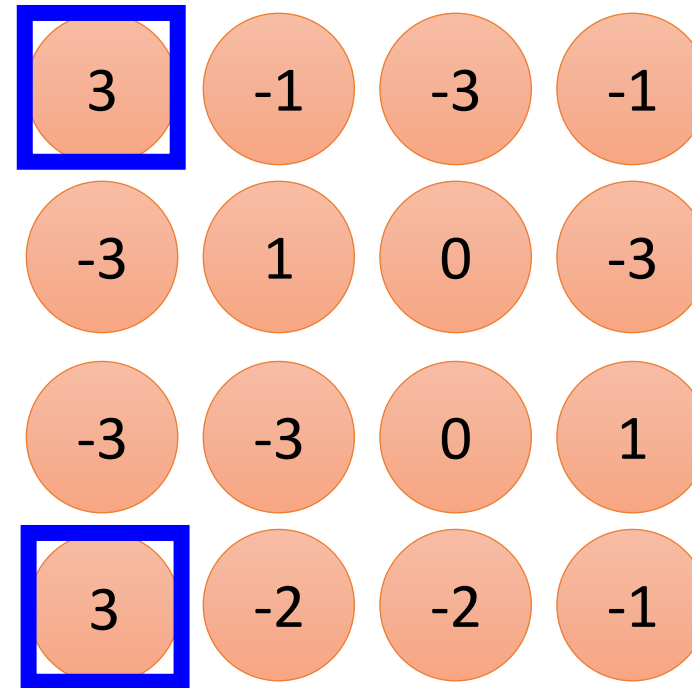
6 x 6 image



A 3x3 grid of numbers representing a filter. The values are: Row 1: 1, -1, -1; Row 2: -1, 1, -1; Row 3: -1, -1, 1. A blue diagonal line runs from the top-left to the bottom-right of the grid.

1	-1	-1
-1	1	-1
-1	-1	1

Filter 1



A 4x4 grid of circles containing numbers. The values are: Row 1: 3, -1, -3, -1; Row 2: -3, 1, 0, -3; Row 3: -3, -3, 0, 1; Row 4: 3, -2, -2, -1. Two blue squares highlight the top-left circle (3) and the bottom-left circle (3).

3	-1	-3	-1
-3	1	0	-3
-3	-3	0	1
3	-2	-2	-1

Property 2

CNN – Convolution

stride=1


1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
1	0	0	0	1	0
0	1	0	0	1	0
0	0	1	0	1	0

6 x 6 image

-1	1	-1
-1	1	-1
-1	1	-1

Filter 2

Do the same process for every filter



4 x 4 image

CNN – Zero Padding

1	-1	-1
-1	1	-1
-1	-1	1

Filter 1

0	0	0				
0	1	0	0	0	0	1
0	0	1	0	0	1	0
	0	0	1	1	0	0
	1	0	0	0	1	0
	0	1	0	0	1	0
	0	0	1	0	1	0
						0
						0
						0


6 x 6 image

You will get another 6 x 6 images in this way




Zero padding


Convolution v.s. Fully Connected




Fully-
connected


1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
1	0	0	0	1	0
0	1	0	0	1	0
0	0	1	0	1	0






Less parameters!



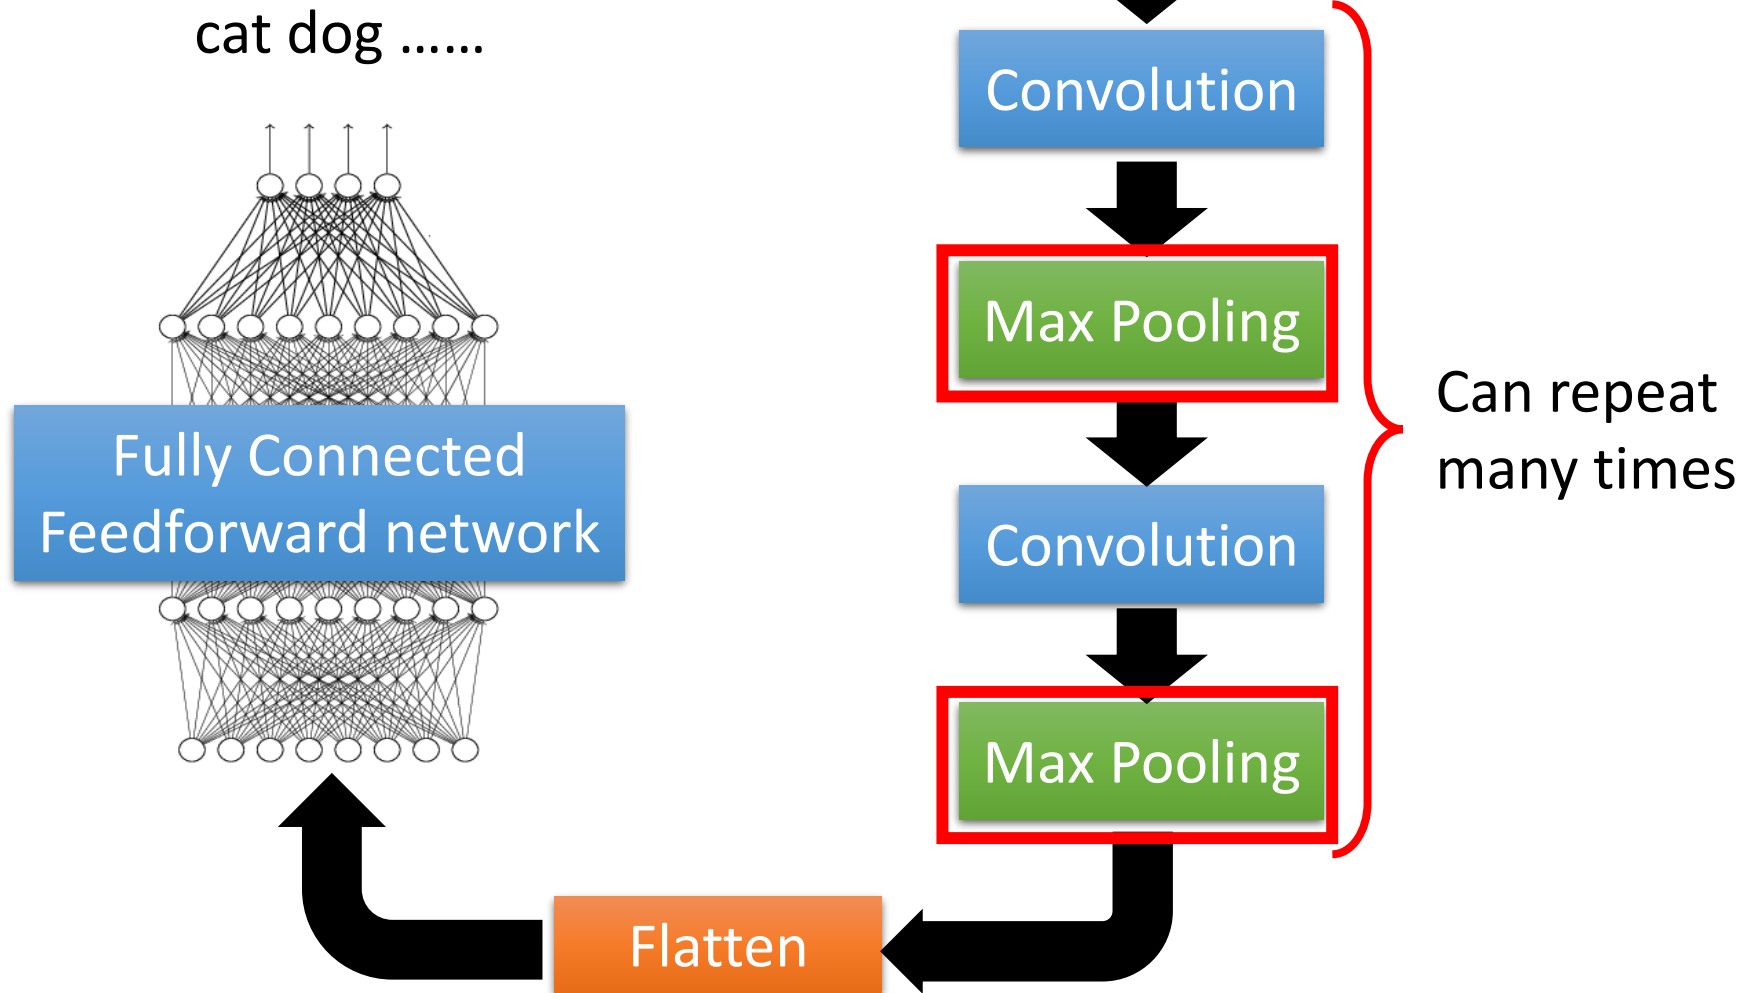


Less parameters!

Even less parameters!



The whole CNN



CNN – Max Pooling

1	-1	-1
-1	1	-1
-1	-1	1

Filter 1

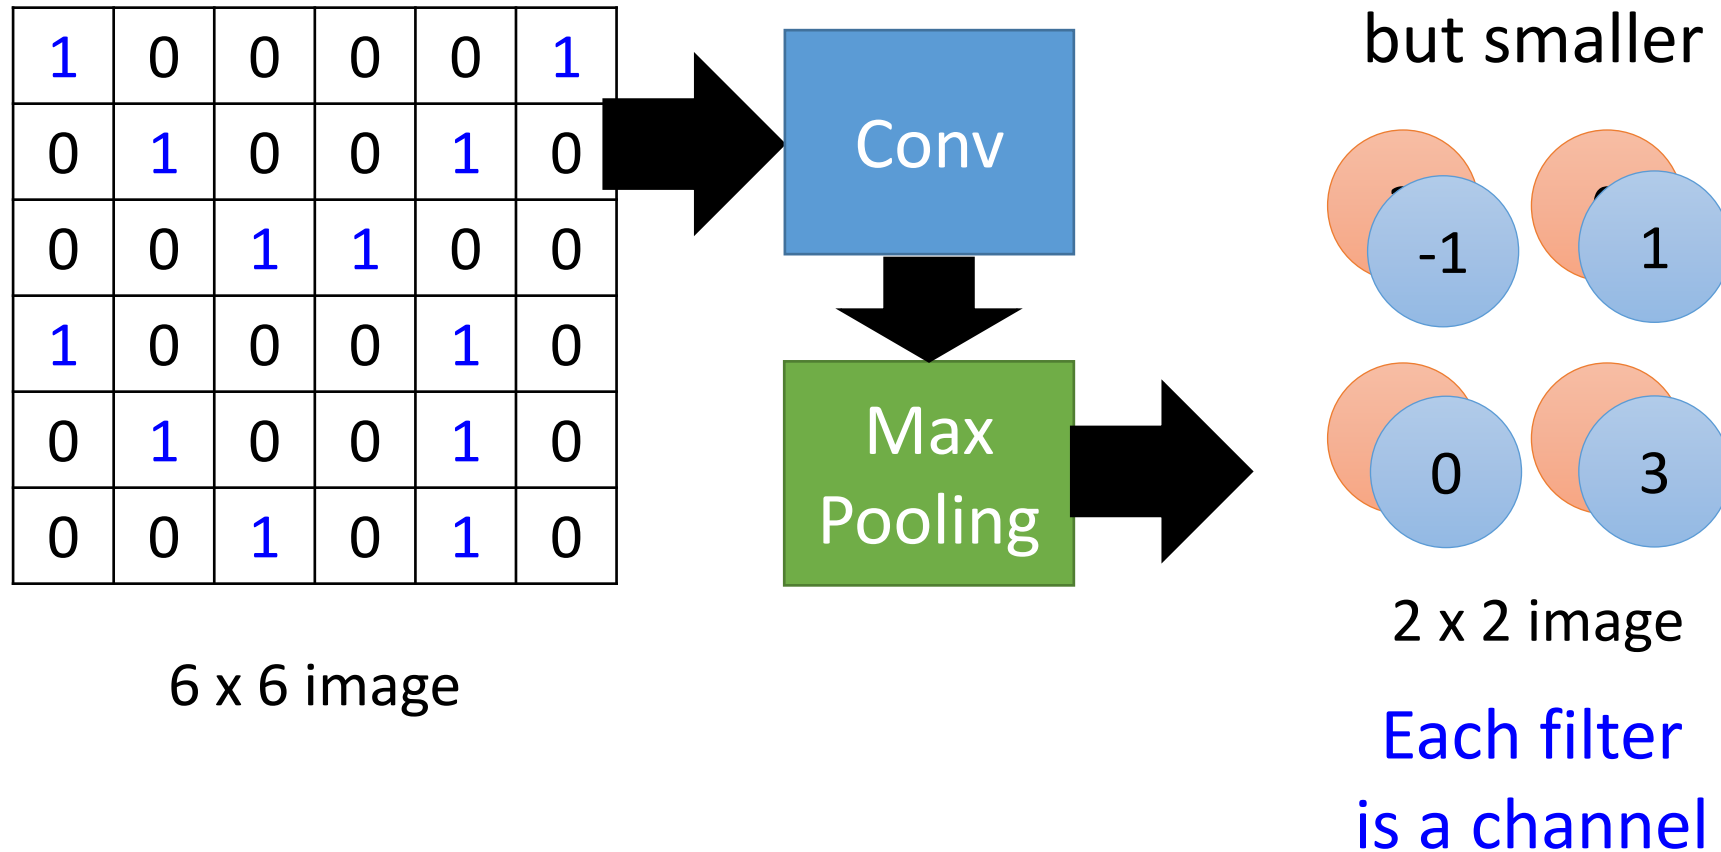
-1	1	-1
-1	1	-1
-1	1	-1

Filter 2


3	-1	-3	-1
-3	1	0	-3
-3	-3	0	1
3	-2	-2	-1

-1	-1	-1	-1
-1	-1	-2	1
-1	-1	-2	1
-1	0	-4	3

CNN – Max Pooling




The whole CNN



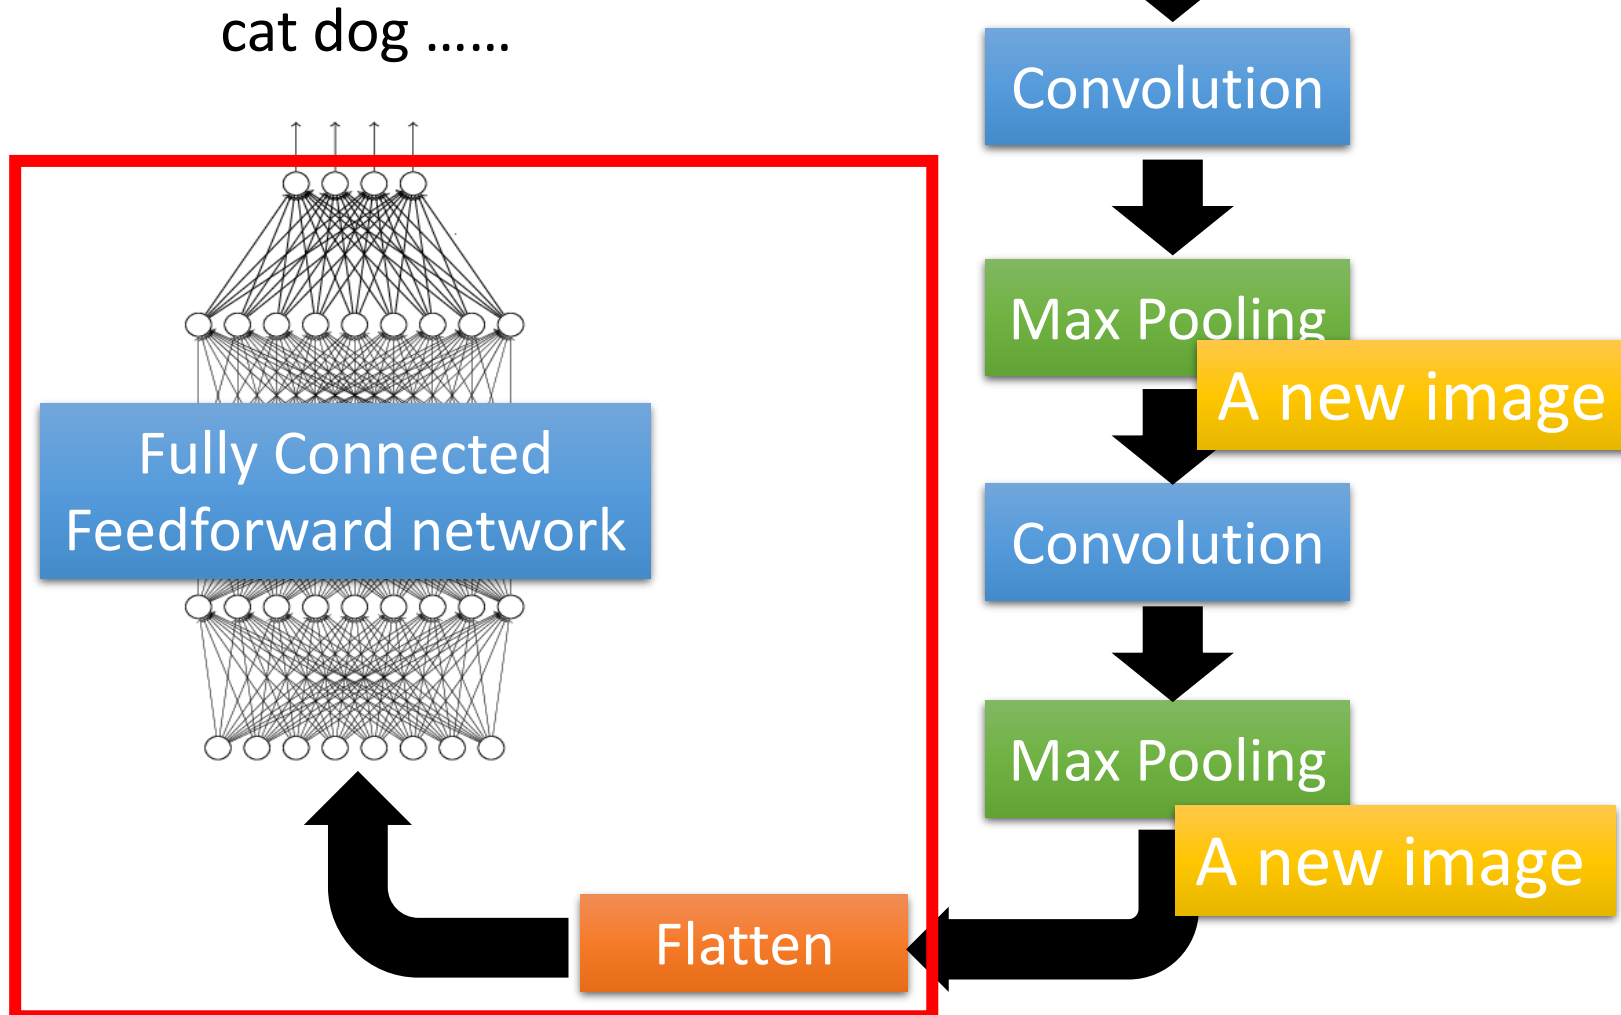
Smaller than the original image

The number of the channel is the number of filters

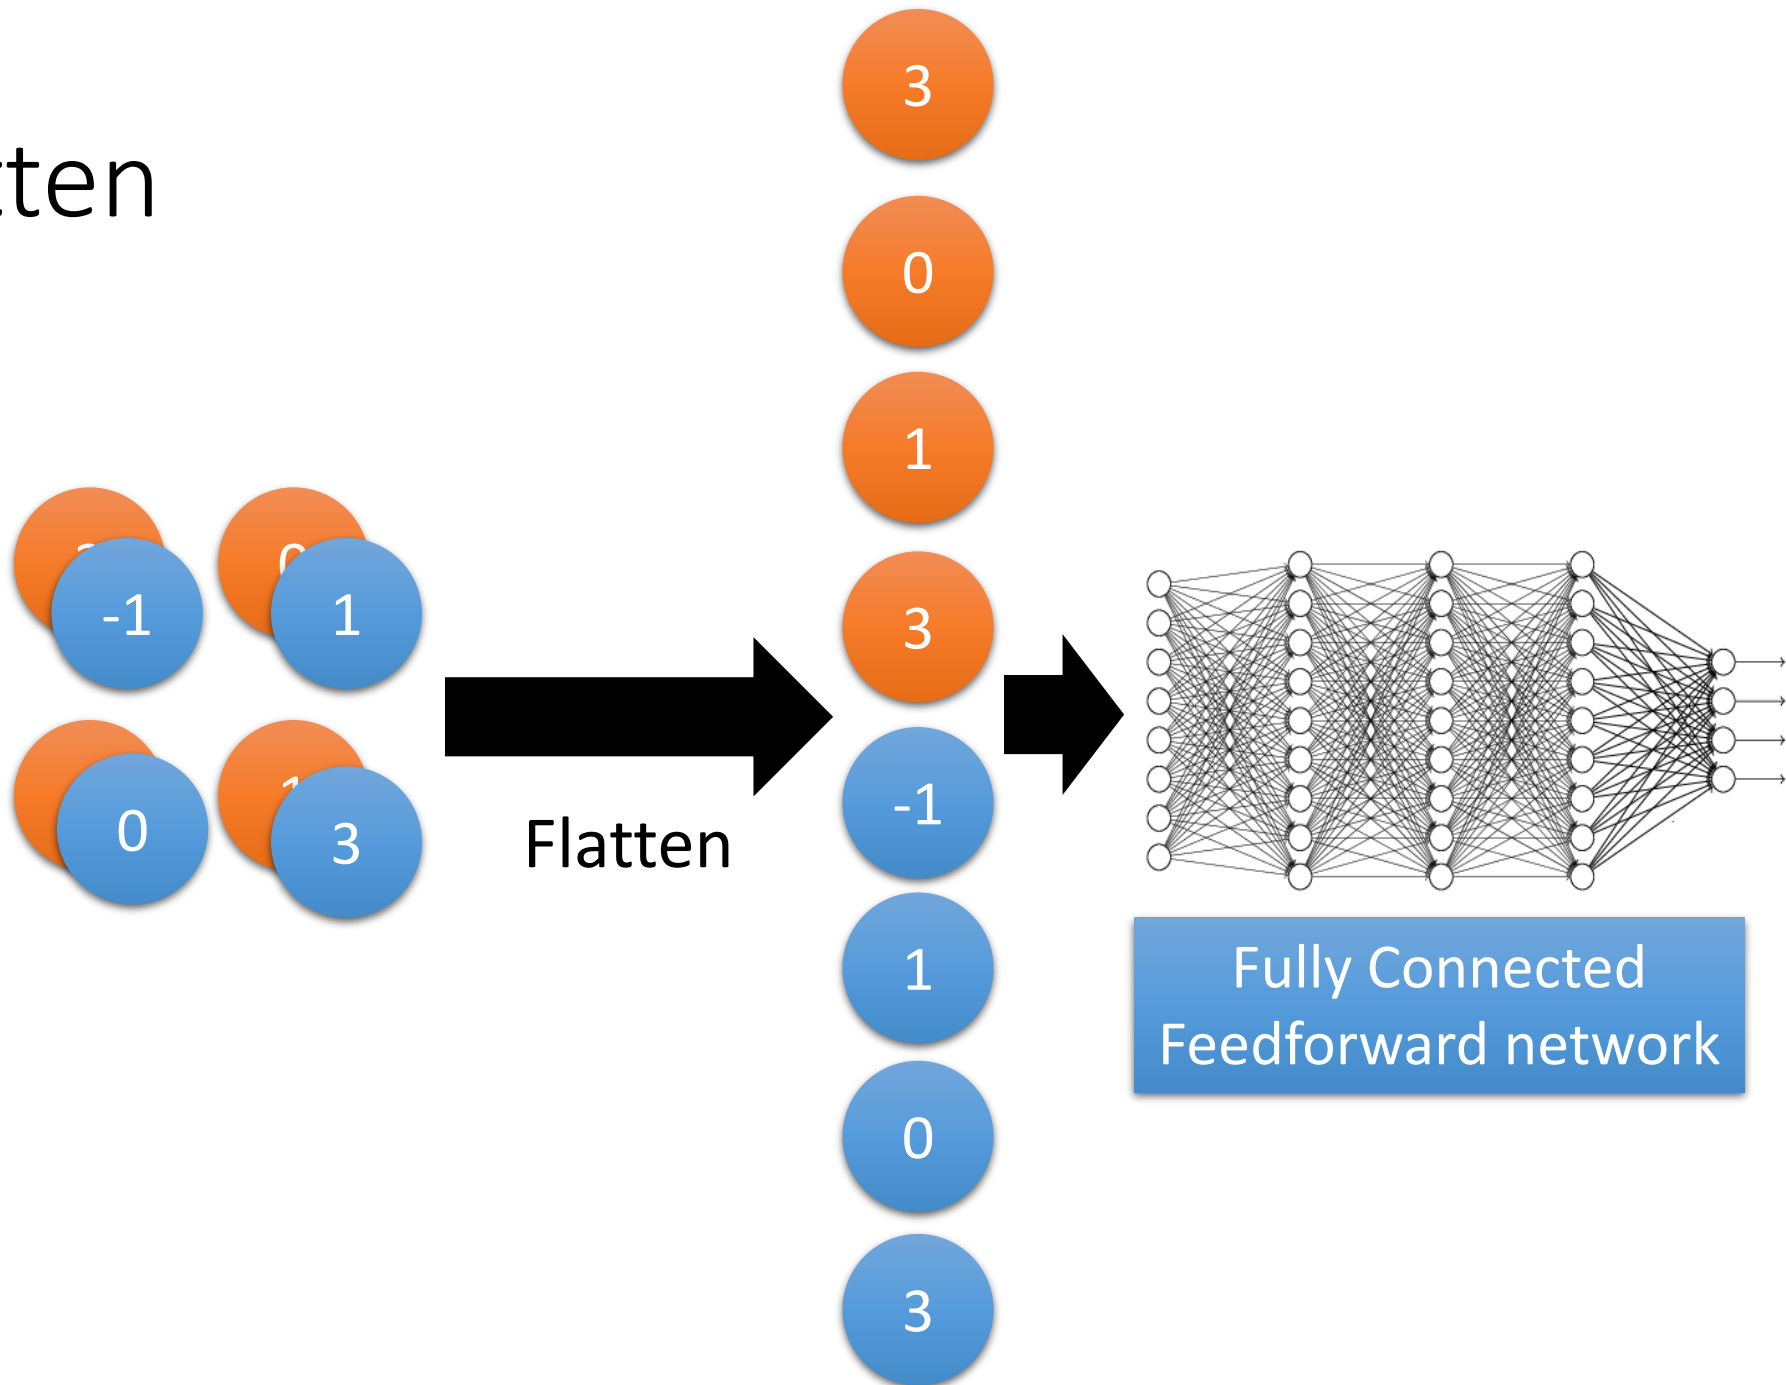


Can repeat many times


The whole CNN



Flatten



Convolutional Neural Network




Learning: Nothing special, just gradient descent

Dealing with RGB Images

Convolution Layer

32x32x3 image




5x5x3 filter



Convolve the filter with the image
i.e. “slide over the image spatially,
computing dot products”

Convolution Layer

32x32x3 image



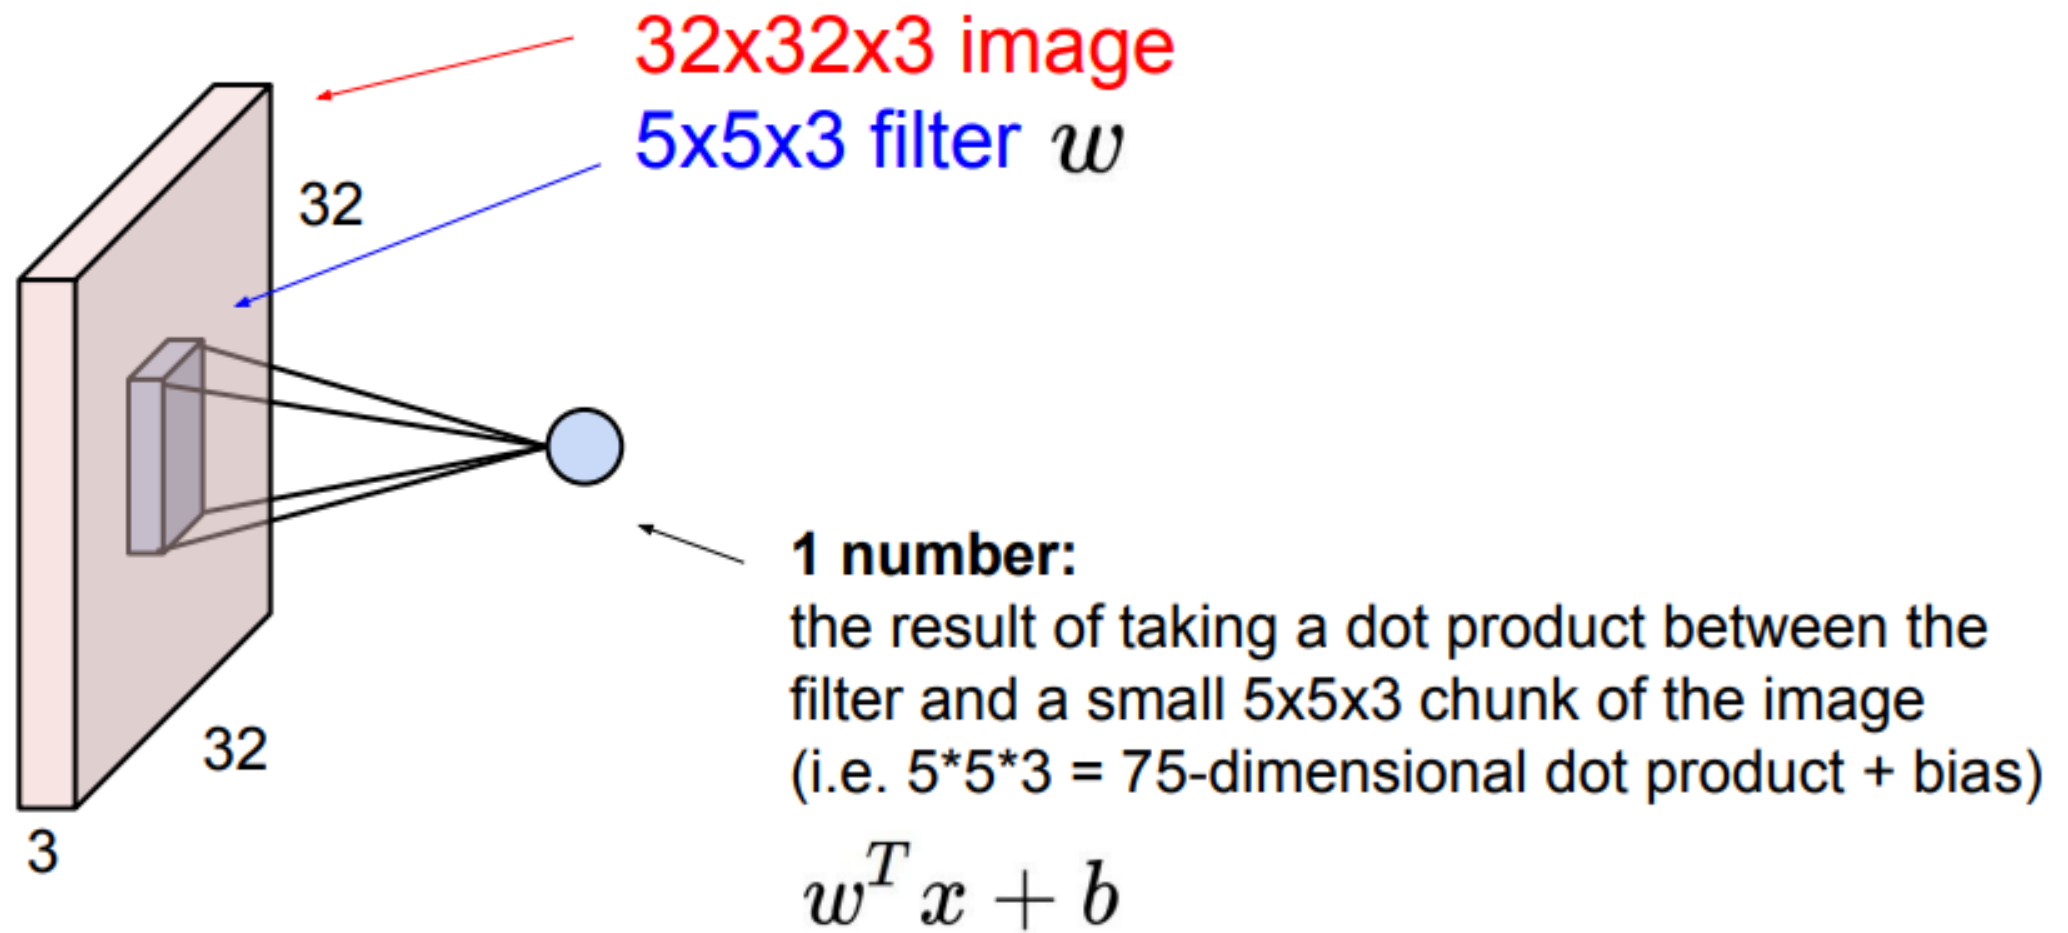
Filters always extend the full depth of the input volume

5x5x3 filter




Convolve the filter with the image
i.e. “slide over the image spatially,
computing dot products”

Convolution Layer




Convolution Layer




Convolution Layer

consider a second, **green** filter




For example, if we had 6 5x5 filters, we'll get 6 separate activation maps:



We stack these up to get a “new image” of size 28x28x6!

Preview: ConvNet is a sequence of Convolutional Layers, interspersed with activation functions



Preview

[Zeiler and Fergus 2013]

Visualization of VGG-16 by Lane McIntosh. VGG-16 architecture from [Simonyan and Zisserman 2014].




Low-level features

Mid-level features

High-level features


Linearly separable classifier



VGG-16 Conv1_1



VGG-16 Conv3_2



VGG-16 Conv5_3

two more layers to go: POOL/FC

