




CSET335L (Deep Learning)

Deep learning

Dr. Rohit Kumar Kaliyar

Outline

- Introduction to Deep Learning (DL)
- Why deep learning?
- Machine learning (ML)
 - Features
 - Weights
 - Artificial Neural Networks (ANN)
 - Loss function
 - Cost function
- History of Deep Learning



What is Deep learning?


Artificial Intelligence

Any technique that enables computers to mimic human behavior




Machine Learning

Ability to program without being explicitly programmed

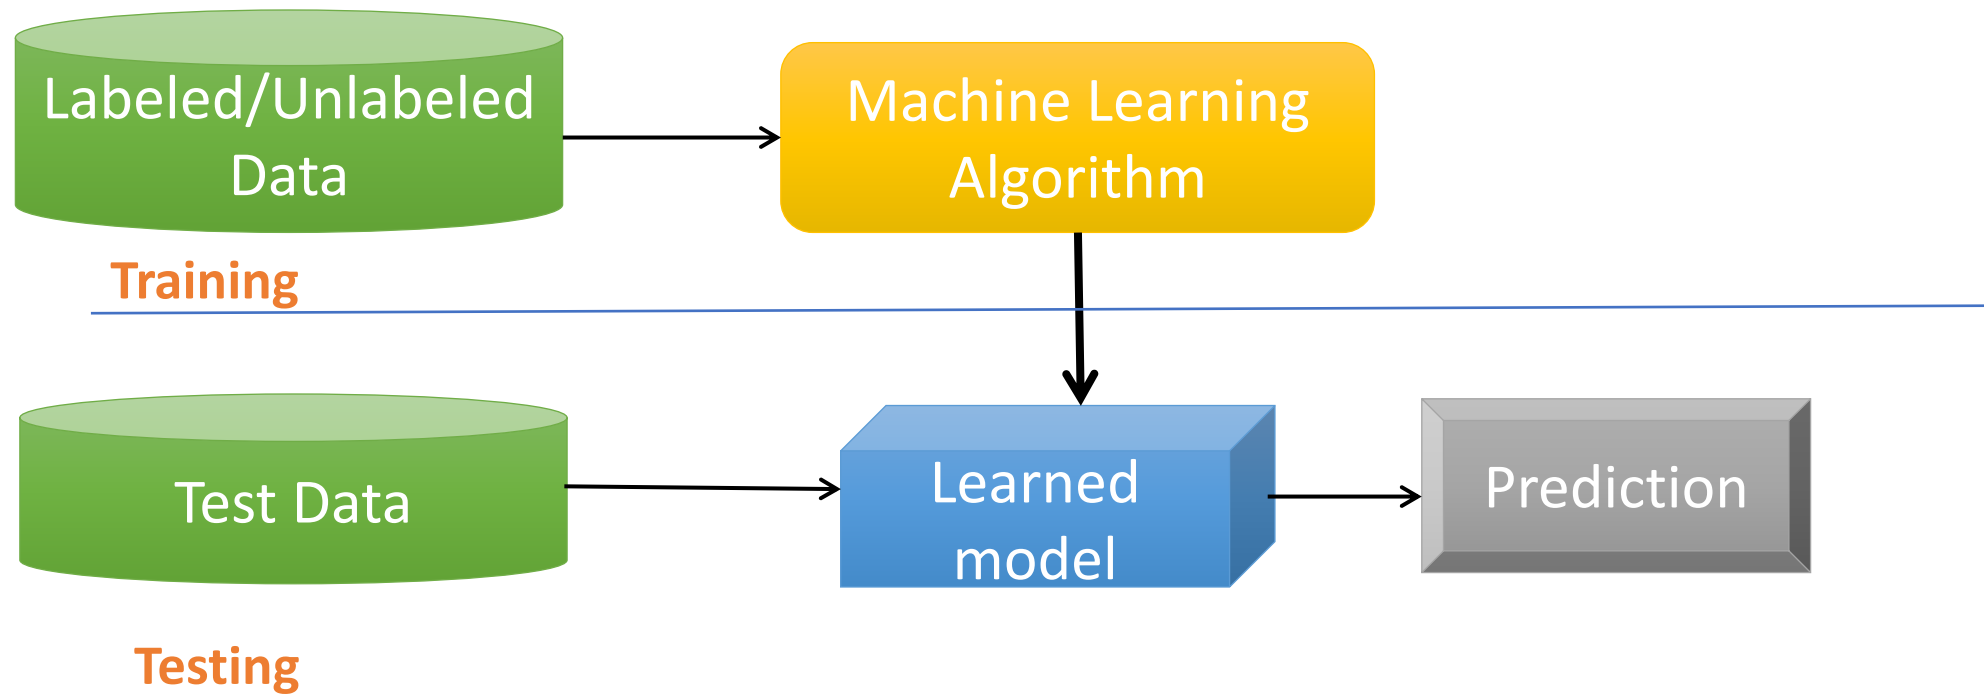



Deep Learning

Extract patterns from data using neural networks




ML gives computers the ability to learn without being explicitly programmed





Low Level Features



Lines & Edges

Mid Level Features



Eyes & Nose & Ears

High Level Features



Facial Structure

Learn

- Deep learning is a type of machine learning that uses multiple layers of artificial neural networks to learn hierarchical features from data. It can apply machine learning to take decisions.

A subfield of machine learning where data representation is learned


Deep learning algorithms attempt to learn (multiple levels of) representation by using a hierarchy of multiple layers

Require tons of data before it begins to understand it and respond in useful ways.


Manually designed features are often over-specified, incomplete, time-consuming, and non-scalable.

Deep learning learns the features directly from data

Machine Learning



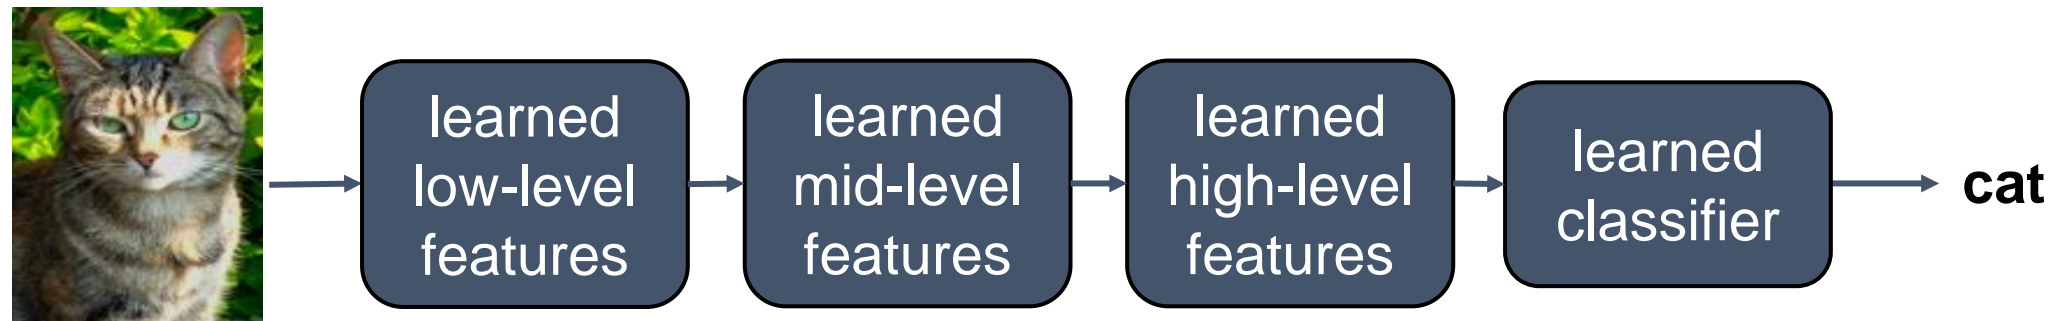
Deep Learning




“Traditional” machine learning:



Deep, “end-to-end” learning:





Neural networks (NN) were first proposed in 1940's and many different architectures of NN have been proposed since then,

So why is the sudden explosion of use of Deep learning now?

Big Data



Hardware
GPUs



Software
New models
and toolboxes






Classification

Classification + Localization





Object Detection




**Speech Recognition
Voice assistant**



Self-driving cars

Perceptron (Single layer neural network) is the structural building block of deep learning models



Inputs Weights Sum Non-Linearity Output

Output


Linear combination of inputs

$$\hat{y} = g \left(w_0 + \sum_{i=1}^m x_i w_i \right)$$


Non-linear activation function


Bias

NN



Neuron forward pass






Inputs Weights Sum Non-Linearity Output

$$\hat{y} = g(w_0 + \mathbf{X}^T \mathbf{W})$$


$$\text{where: } \mathbf{X} = \begin{bmatrix} x_1 \\ \vdots \\ x_m \end{bmatrix} \text{ and } \mathbf{W} = \begin{bmatrix} w_1 \\ \vdots \\ w_m \end{bmatrix}$$

- A neural network is comprised of layers of nodes and learns to map examples of inputs to outputs.
- For a given node, the inputs are multiplied by the weights in a node and summed together. This value is then transformed via an activation function and defines the specific output or “activation” of the node.
- The simplest activation function is referred to as **linear activation**, where no transform is applied at all.
- A network comprised of only linear activation functions is very easy to train but cannot learn complex mapping functions. Linear activation functions are still used in the output layer for networks that predict a quantity (e.g., regression problems).
- **Nonlinear activation functions** are preferred as they allow the nodes to learn more complex structures in the data.

- Activation function introduce non-linearity in the network
- Non-linear functions approximate complex functions that make neural networks extremely powerful




Linear activation functions produce linear decisions no matter the network size



Non-linearities allow us to approximate arbitrarily complex functions

Example: Suppose we have a trained network with weights W and two inputs x_1 and x_2 .




We have: $w_0 = 1$ and $\mathbf{W} = \begin{bmatrix} 3 \\ -2 \end{bmatrix}$

$$\begin{aligned}\hat{y} &= g(w_0 + \mathbf{X}^T \mathbf{W}) \\ &= g\left(1 + \begin{bmatrix} x_1 \\ x_2 \end{bmatrix}^T \begin{bmatrix} 3 \\ -2 \end{bmatrix}\right) \\ \hat{y} &= g(1 + 3x_1 - 2x_2)\end{aligned}$$

This is just a line in 2D!

Till the summation step, if we feed $x_1 = -1$ and $x_2 = 2$, we get -6.

$$\hat{y} = g(1 + 3x_1 - 2x_2)$$



Assume we have input: $\mathbf{X} = \begin{bmatrix} -1 \\ 2 \end{bmatrix}$

$$\begin{aligned}\hat{y} &= g(1 + (3 * -1) - (2 * 2)) \\ &= g(-6) \approx 0.002\end{aligned}$$

- Unit step function (Output a 0 or 1)
- Sigmoid / Logistic function (S-shaped | Output value between 0 and 1)
- Hyperbolic tangent (tanh) function (Output value between -1 to 1)
- Rectified linear (ReLU) function (Output value between 0 to infinity)

Sigmoid function: Non-zero centered
Vanishing gradient problem

Tanh function: Zero-centered
Vanishing gradient problem

ReLu function: Non-zero centered
Vanishing gradient problem does not exist

More reading on activation functions:


<https://www.analyticsvidhya.com/blog/2021/04/activation-functions-and-their-derivatives-a-quick-complete-guide/>

For a long time till early 1990s, **sigmoid function** was the **default activation** used on neural networks.

In the later 1990s and through the 2000s, the **tanh was preferred over the sigmoid** function as it makes training easier and had better predictive performance.

Later, **Relu is mostly used for deep learning models** because of its ability to map complex relationships


$$\hat{y} = g(w_0 + \mathbf{X}^T \mathbf{W})$$



3 steps of Perceptron:


1. Dot product (Inputs multiplied with weights)
2. Add bias
3. Apply non-linearity

Here all inputs are connected to all outputs, these layers are called dense layers



$$z_i = w_{0,i} + \sum_{j=1}^m x_j w_{j,i}$$

Here all inputs are connected to all outputs, these layers are called dense layers



x_1

x_2

x_m

Inputs



z_1

z_2

z_3

z_n

Hidden



\hat{y}_1


\hat{y}_2

Output



```
import tensorflow as tf
```

```
model = tf.keras.Sequential([  
    tf.keras.layers.Dense(n),  
    tf.keras.layers.Dense(2)  
])
```

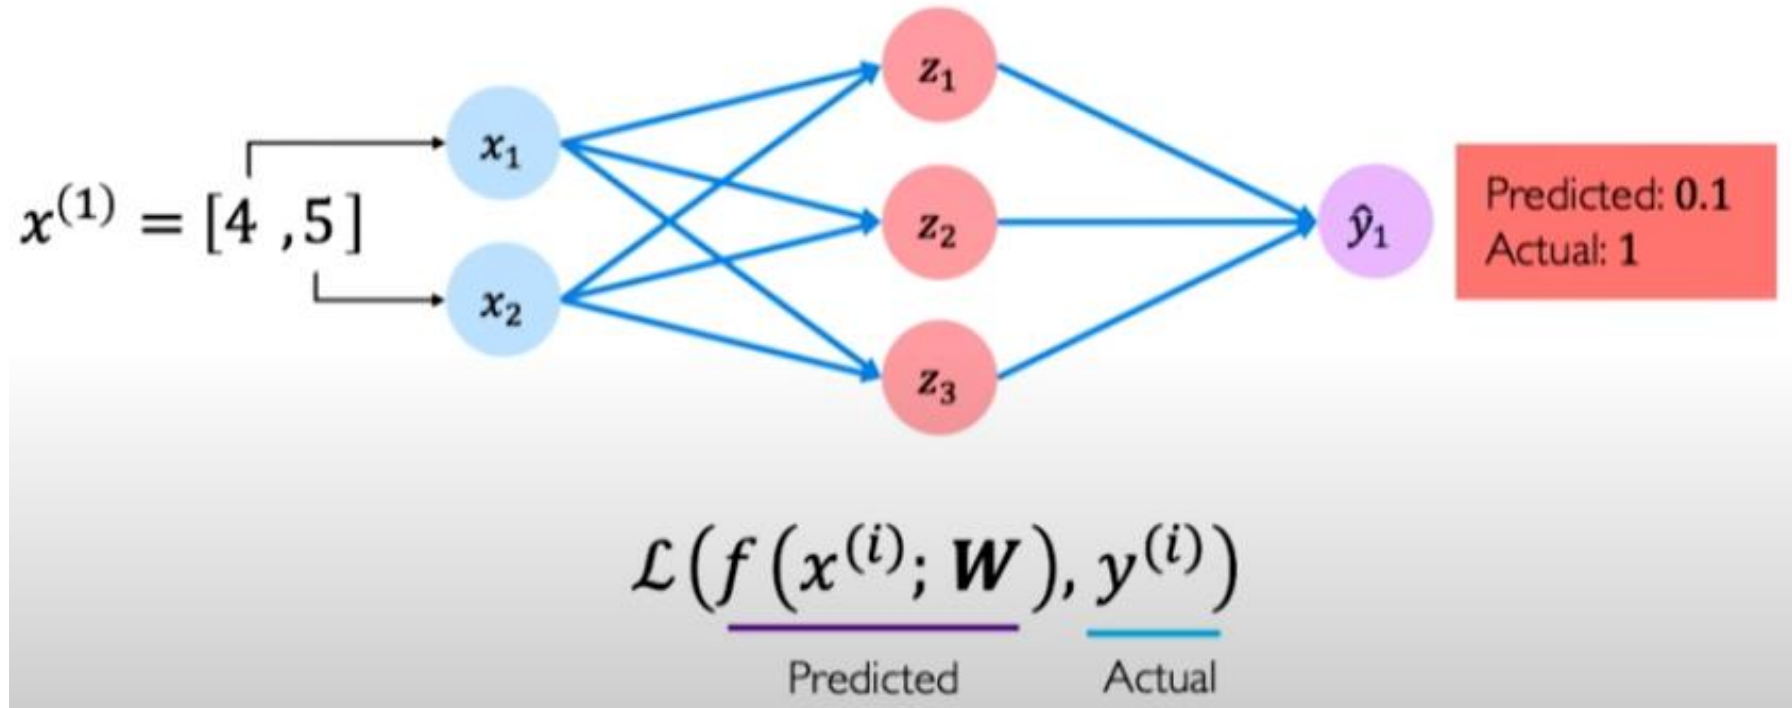



```
import tensorflow as tf

model = tf.keras.Sequential([
    tf.keras.layers.Dense(n_1),
    tf.keras.layers.Dense(n_2),
    :
    tf.keras.layers.Dense(2)
])
```

$$z_{k,i} = w_{0,i}^{(k)} + \sum_{j=1}^{n_{k-1}} g(z_{k-1,j}) w_{j,i}^{(k)}$$

The loss function measures the cost incurred by incorrect predictions.






Also known as:

- Objective function
- Cost function
- Empirical Risk

$$J(\mathbf{W}) = \frac{1}{n} \sum_{i=1}^n \mathcal{L}(\underbrace{f(x^{(i)}; \mathbf{W})}_{\text{Predicted}}, \underbrace{y^{(i)}}_{\text{Actual}})$$

Output a probability between 0 and 1



$$J(W) = -\frac{1}{n} \sum_{i=1}^n \underbrace{y^{(i)}}_{\text{Actual}} \log \left(\underbrace{f(x^{(i)}; W)}_{\text{Predicted}} \right) + (1 - \underbrace{y^{(i)}}_{\text{Actual}}) \log \left(1 - \underbrace{f(x^{(i)}; W)}_{\text{Predicted}} \right)$$

Used for regression models; Output continuous numbers

